



Upgrade to paid



data model predominates. Output primarily depends on what was fed into the language model, not on whether it is cohesive.

Now, as it turns out, language by itself does have a certain innate cohesion, but it's relatively weak; this is one of the reasons that you can apply a data model to an LLM and get something that seems structured (it may be garbage data, but at least it's in a workable format). However, one key benefit of a knowledge graph is that, if designed well, it can produce a cohesive, structured narrative. A typical knowledge graph is much (several orders of magnitude) smaller than the equivalent LLM, so this obviously has environmental and cost benefits as well.

When you willy-nilly drop in different data sources without attempting to create a canonical model, what you end up with is difficult to query, is subject to other rules and constraints depending upon the data set, and may just be poorly or incorrectly modelled. We are reaching a stage with generative AI where it is possible to create syncretic solutions at relatively minimal token costs, primarily by controlling the transformations of data sources (various heresies) into an established orthodoxy that varies from one organisational entity to the next.

Converting data directly with LLMs doesn't scale well (it requires a significant portion of the model to be maintained in context). Still, if you have context-aware transformations (such

